

Dylan Hutchison

📞 862 226 2764 • ✉ dhutchis@cs.washington.edu
🌐 [linkedin.com/in/dylanhutchison](https://www.linkedin.com/in/dylanhutchison) • github.com/dhutchis

Objective: To bridge linear and relational algebra systems both in theory and in practice.
To build the *polystore optimizer* that translates and optimizes queries across systems, taking all known statistics, properties, and equivalences into account.

Education

University of Washington

Seattle, WA

Ph.D. in Computer Science & Engineering, M.S. in Computer Science & Engineering (3/2017)

9/2015–≈2019

Awards NSF Graduate Research Fellow

Advisors Bill Howe, Dan Suciu, Zachary Tatlock

Stevens Institute of Technology

Hoboken, NJ

M.S. in Computer Science, M.S. in Applied Mathematics, B.E. in Computer Engineering

8/2010–5/2015

GPA 4.00 Graduate, 3.97 Undergraduate

Thesis *ModelWizard: Toward Interactive Model Construction*

advised by David A. Naumann, Philippos Mordohai, Andrew D. Gordon

Awards 2014 National Barry Goldwater Scholar, Association of Old Crows Scholar, Tau Beta Pi Scholar, Computing Research Association *Outstanding Undergraduate Researcher* Honorable Mention

Societies *Tau Beta Pi* (Engineering), *Upsilon Pi Epsilon* (Computer Science), *Eta Kappa Nu* (IEEE)

University of Edinburgh

Edinburgh, UK

Study Abroad Semester, 6 courses transferred

1/2014–5/2014

Experience

Apache Accumulo – PMC Committer

10/2015–Future

- Contribute code and discussion to the Apache Accumulo database community

Laboratory & Industry

IBM Research Almaden – Research Intern

San Jose, CA

SystemML Team; Mentors Matthias Boehm, Alexandre Evfimievski

6/2017–8/2017

- Develop cost-based optimization for Apache SystemML based on elementary rewrites over sum-product operators
- Achieve speedups as large as 20,000x on machine learning algorithm kernels

MIT Lincoln Laboratory – Research Engineer

Lexington, MA

Computing and Analytics Group; Advisors Jeremy Kepner, Vijay Gadepally

1/2015–9/2015

- Engineered Graphulo, a Java server-side matrix math library for the Accumulo database
- Recasted graph algorithms into the GraphBLAS standard; prototyped in Matlab

Microsoft Research – Research Intern

Cambridge, UK

Programming Principles and Tools Group; Advisor Andy Gordon

6/2014–8/2014

- Designed ModelWizard: a DSL in F# for interactive model construction targeting Tabular, a schema-based probabilistic programming language. Presented a concept poster at the Microsoft PhD Summer School

Sandia National Laboratories – Technical Intern

Livermore, CA

Information Assurance Group; Advisors Levi Lloyd, Tamara Kolda

5/2013–8/2013

- Pursued network anomaly detection via Accumulo schemas, machine learning and visualization
- Scaled LXC (Linux Containers) with MiniMega, a mass distributed VM experiment platform

MIT Lincoln Laboratory – Research Intern

Lexington, MA

Computing and Analytics Group; Advisor Jeremy Kepner

5/2012–8/2012

Bioengineering and Systems Technology Group; Advisor Darrell Ricke

- Integrated and benchmarked Accumulo distributed database features into D4M, a Matlab package delivering linear algebra and graph theory capabilities via Associative Arrays
- Applied D4M work to a DNA matching bioinformatics project, published in the *Lincoln Laboratory Journal*

Brown Brothers Harriman – Web Development Co-op

Business Application Development; Advisors John David, Steve Hansen

Jersey City, NJ

1/2012–5/2012

- o Designed and developed front- and back-end web applications for financial reporting using SQL, C++, and jQuery

Teaching.....

Stevens Institute of Technology

Hoboken, NJ

Computer Science Department: Teaching Assistant

8/2012–12/2013

- o Teach, create and evaluate computer science coursework for classes ranging from 40 up to 70 students
- o CS 506: Intro to IT Security, CS 135: Discrete Structures, CS 334: Automata and Computation

Academic Support Center: Tutor

8/2011–12/2013

- o Teach individuals and groups in Mathematics, Computer Science, and Engineering

Activities

Global Grand Challenges Summit: 2017, authored a podcast with Sarah Engel on *The “Power” of the Internet*

▷ <https://www.ucl.ac.uk/steapp/professional-education/ggcs-how-to-change-the-world/podcasts/sustainability/sustainability-6>

Reviewer for Information Systems: 2016–2017, a journal published by Wiley

Myria Database Maintainer: Maintain the Myria database alongside UWDB students

Undergrad Activities: Organized seminars as VP of the *Graduate Computer Science Society*; hosted a road race as President of the *Cycling Club*; led operations for the 2012 *Castle Point Anime Convention*, attracting over 2100 people

Publications

Conference Papers.....

- [C12] D. Hutchison, **Distributed triangle counting in the Graphulo matrix math library**, in *High Performance Extreme Computing (HPEC)*, IEEE, Sep. 2017, Graph Challenge Honorable Mention.
- [C11] L. Milechin, A. Chen, V. Gadepally, D. Hutchison, S. Samsi, and J. Kepner, **D4M 3.0: Extended database and language capabilities**, in *High Performance Extreme Computing (HPEC)*, IEEE, Sep. 2017. arXiv: 1702.03253 [cs.DB].
- [C10] D. Hutchison, B. Howe, and D. Suciu, **LaraDB: A minimalist kernel for linear and relational algebra computation**, in *SIGMOD Workshop on Algorithms and Systems for MapReduce and Beyond (BeyondMR)*, ACM, May 2017. DOI: 10.1145/3070607.3070608. arXiv: 1703.07342 [cs.DB]. Online: <https://youtu.be/d-ZY8lIs5Pc?t=2m45s>, Slides: <https://sites.google.com/site/beyondmr2017/program>.
- [C9] J. Wang, T. Baker, M. Balazinska, D. Halperin, B. Haynes, B. Howe, D. Hutchison, S. Jain, R. Maas, P. Mehta, D. Moritz, B. Myers, J. Ortiz, D. Suciu, A. Whitaker, and S. Xu, **The Myria big data management and analytics system and cloud service**, in *Conference on Innovative Data Systems Research (CIDR)*, Jan. 2017. Online: <https://homes.cs.washington.edu/~magda/papers/wang-cidr17.pdf>, Slides: <http://cidrdb.org/cidr2017/slides/p37-wang-cidr17-slides.pdf>.
- [C8] D. Hutchison, J. Kepner, V. Gadepally, and B. Howe, **From NoSQL Accumulo to NewSQL Graphulo: Design and utility of graph algorithms inside a BigTable database**, in *High Performance Extreme Computing (HPEC)*, **Best Student Paper**, IEEE, Sep. 2016. DOI: 10.1109/HPEC.2016.7761577. arXiv: 1606.07085 [cs.DB], Slides: <https://github.com/Accla/graphulo/blob/master/docs/presentations/2016-09-HPEC-Graphulo-Algorithms.pdf>.
- [C7] J. Kepner, V. Gadepally, D. Hutchison, H. Jananthan, T. Mattson, S. Samsi, and A. Reuther, **Associative array model of SQL, NoSQL, and NewSQL databases**, in *High Performance Extreme Computing (HPEC)*, IEEE, Sep. 2016. DOI: 10.1109/HPEC.2016.7761647. arXiv: 1606.05797 [cs.DB].
- [C6] T. Weale, V. Gadepally, D. Hutchison, and J. Kepner, **Benchmarking the Graphulo processing framework**, in *High Performance Extreme Computing (HPEC)*, IEEE, Sep. 2016. DOI: 10.1109/HPEC.2016.7761640. arXiv: 1609.08642 [cs.DB].

- [C5] A. Chen, A. Edelman, J. Kepner, V. Gadepally, and D. Hutchison, **Julia implementation of the dynamic distributed dimensional data model**, in *High Performance Extreme Computing (HPEC)*, **Best Paper**, IEEE, Sep. 2016. DOI: 10.1109/HPEC.2016.7761626. arXiv: 1608.04041 [cs.MS].
- [C4] J. Kepner, P. Aaltonen, D. Bader, A. Buluç, F. Franchetti, J. Gilbert, D. Hutchison, M. Kumar, A. Lumsdaine, H. Meyerhenke, S. McMillan, J. Moreira, J. D. Owens, C. Yang, M. Zalewski, and T. Mattson, **Mathematical foundations of the GraphBLAS**, in *High Performance Extreme Computing (HPEC)*, IEEE, Sep. 2016. DOI: 10.1109/HPEC.2016.7761646. arXiv: 1606.05790 [cs.MS].
- [C3] D. Hutchison, J. Kepner, V. Gadepally, and A. Fuchs, **Graphulo implementation of server-side sparse matrix multiply in the Accumulo database**, in *High Performance Extreme Computing (HPEC)*, **Best Student Paper Finalist**, IEEE, Sep. 2015. DOI: 10.1109/HPEC.2015.7322448. arXiv: 1507.01066 [cs.DB], Slides: <https://github.com/Accla/graphulo/blob/master/docs/presentations/2015-09-HPEC-Graphulo-MatrixMultiply.pdf>.
- [C2] V. Gadepally, J. Bolewski, D. Hook, D. Hutchison, B. Miller, and J. Kepner, **Graphulo: Linear algebra graph kernels for NoSQL databases**, in *International Parallel & Distributed Processing Symposium Workshops (IPDPSW)*, IEEE, May 2015. DOI: 10.1109/IPDPSW.2015.19. arXiv: 1508.07372 [cs.DS], Slides: <https://github.com/Accla/graphulo/blob/master/docs/presentations/2015-05-GABB-GraphuloInGraphBLAS.pdf>.
- [C1] D. Hutchison and S. Kleinberg, **Causal inference under uncertainty via adjustments and SOPDs**, in *Causality and Experimentation in the Sciences*, Paris, France, Jul. 2013. Online: http://caeits.sciencesconf.org/conference/caeits/hutchison_caeits2013.pdf, Slides: <https://github.com/dhutchis/documents/blob/master/presentations/2013-06-caeits-sopd.pdf>.

Journal Papers.....

- [J2] M. J. Smith, V. Vaglica, M. Sajeve, N. McGough, D. Hutchison, A. D. Gordon, C. Russo, A. Ramarosandratana, and W. Stuppy, **Monitoring internet trade to inform species conservation actions**, *Endangered Species Research*, Mar. 2017. DOI: 10.3354/esr00803.
- [J1] J. Kepner, D. Ricke, and D. Hutchison, **Taming biological big data with D4M**, *Lincoln Laboratory Journal*, vol. 20, no. 1, 2013. Online: https://www.ll.mit.edu/publications/journal/pdf/vol20_no1/20_1_6_Kepner.pdf.

Thesis.....

- [T1] D. Hutchison, **ModelWizard: Toward interactive model construction**, M.S. Thesis, Stevens Institute of Technology, May 2015. arXiv: 1604.04639 [cs.PL].

Preprints.....

- [U1] D. Hutchison, B. Howe, and D. Suciu, **Lara: A key-value algebra underlying arrays and relations**, Apr. 2016. arXiv: 1604.03607 [cs.DB].

Blog Posts.....

- [B1] D. Hutchison, B. Howe, D. Suciu, and Z. Tatlock. (Apr. 2016). **PolyPEG: A proposal for polystore optimization**, Online: <http://istc-bigdata.org/index.php/polypeg-a-proposal-for-polystore-optimization>.

Posters.....

- [P10] D. Hutchison, M. Boehm, A. Evfimievski, B. Reinwald, P. Sen, and F. Ozcan, **Holistic sum-product optimization for large-scale machine learning**, Aug. 2017, Poster presented at IBM Intern Symposium.
- [P9] D. Hutchison, B. Howe, V. Gadepally, and J. Kepner, **In-database vs. external system analytics on a key-value store**, Jan. 2017. Online: <https://github.com/Accla/graphulo/blob/master/docs/posters/2017-01-Graphulo-MapReduce.pdf>, Poster presented at North East Database Day (NEDB).

- [P8] D. Hutchison, S. Jain, B. Howe, and D. Maier, **Ocean genomic analysis with Myria**, Aug. 2016. Online: <https://github.com/uwdb/related-work/blob/master/oceanography-istc/2016-08-Myria-ocean-ISTC.pdf>, Poster presented at summit for the Intel Science and Technology Center (ISTC) for Big Data.
- [P7] D. Hutchison, V. Gadepally, J. Kepner, and B. Howe, **Graphulo: Native linear algebra in a NoSQL DB**, Aug. 2016. Online: <https://github.com/Accla/graphulo/blob/master/docs/posters/2016-08-Graphulo-algorithms-ISTC.pdf>, Poster presented at summit for the Intel Science and Technology Center (ISTC) for Big Data.
- [P6] D. Hutchison, B. Howe, D. Suciu, and Z. Tatlock, **Polystore optimization via program expression graphs**, Jan. 2016. Online: <https://github.com/dhutchis/documents/blob/master/posters/2016-01-21-polypeg-poster.pdf>, Poster presented at North East Database Day (NEDB).
- [P5] D. Hutchison, J. Kepner, and V. Gadepally, **Graphulo: Graph processing for Accumulo databases**, Jan. 2016. Online: <https://github.com/Accla/graphulo/blob/master/docs/posters/2015-09-Graphulo-ISTC-Poster.pdf>, Poster presented at North East Database Day (NEDB) and Sep. 2015 retreat for the Intel Science and Technology Center (ISTC) for Big Data.
- [P4] L. Battle, L. Edwards, V. Gadepally, B. Gavin, B. Hancock, D. Hutchison, J. Kepner, and A. Moran, **Technologies for visualization of big medical text data**, Aug. 2015. Online: <https://github.com/Accla/graphulo/blob/master/docs/posters/2015-08-BigDAWG-MIMIC-Topic-Modeling.pdf>, Poster presented at Very Large Databases Conference (VLDB).
- [P3] D. Hutchison, E. Cherin, X. Li, and H. Yang, **HBaaS: Heterogeneously-accelerated bioinformatics-as-a-service**, Apr. 2015, Poster and demo presented at Stevens Institute Senior Design Expo.
- [P2] D. Hutchison and the Microsoft Research Tabular Team, **Structural clustering**, Jul. 2014. Online: <https://github.com/dhutchis/documents/blob/master/posters/2014-07-structural-clustering.pdf>, Poster presented at Microsoft PhD Summer School.
- [P1] D. Hutchison and D. Kleszyk, **Investigating the influence of infrastructure on the load response of stevens course servers**, Dec. 2011, Course project for CPE 345 Modeling and Simulation.

Other Talks and Demos.....

- [O11] D. Hutchison, **Holistic sum-product optimization for large-scale machine learning**, Aug. 2017, IBM Project Presentation.
- [O10] —, **In-database analytics for NoSQL key-value stores**, Dec. 2016. Online: <https://www.cs.washington.edu/qualexam/dhutchis>, Slides: <https://github.com/Accla/graphulo/blob/master/docs/presentations/2016-12-Key-Value-Analytics.pdf>.
- [O9] B. Howe, D. Hutchison, and S. Jain, **Demo of myria as a federated database**, Oct. 2016. Online: https://github.com/uwescience/raco/blob/SPJA_federation/HPDA_review.ipynb, Presentation and demo to government sponsors.
- [O8] D. Hutchison, J. Kepner, and V. Gadepally, **Lara: A language of linear and relational algebra for polystores**, Dec. 2015. Online: http://db.cs.washington.edu/events/database_day/2015/database_day_2015.html, UWDB Database Day industry talk.
- [O7] —, **Graphulo use and design**, Aug. 2015, Presentation and demo to government sponsors.
- [O6] V. Gadepally, L. Edwards, D. Hutchison, and J. Kepner, **Using d4m for rapid prototyping of analytics for apache accumulo**, Apr. 2015. Online: <http://accumulosummit.com/program/talks/using-d4m-for-rapid-prototyping-of-analytics-for-apache-accumulo/>, Presentation at Accumulo Summit.
- [O5] D. Hutchison, A. Gordon, and C. Russo, **ModelWizard: Interactive model construction for Tabular**, Aug. 2014. Online: <https://github.com/dhutchis/documents/blob/master/presentations/2014-08-modelwizard-tabular.pdf>, Presentation and demo to Programming Principles and Tools at Microsoft Research.

- [O4] D. Hutchison and L. Lloyd, **Network anomaly detection**, Aug. 2013. Online: <https://github.com/dhutchis/documents/blob/master/presentations/2013-08-network-anomaly-detection.pdf>, Presentation at an intern symposium at Sandia National Laboratories.
- [O3] D. Hutchison, **Our aims as modelers: Toward better predictions, explanations, interventions**, Aug. 2013. Online: github.com/dhutchis/ModelTalk, Presentation at a technical seminar at Sandia National Laboratories.
- [O2] —, **Uncertain? Represent your belief with SOPD**, Jan. 2013. Online: <https://github.com/dhutchis/documents/blob/master/presentations/2013-01-omm-sopd.pdf>, Minute madness talk at Programming Languages Mentoring Workshop (PLMW) at the Principles of Programming Languages Conference (POPL).
- [O1] —, **Accelerating bioinformatics with big data technologies**, Aug. 2012. Online: <https://github.com/dhutchis/documents/blob/master/presentations/2012-08-d4m-bioinformatics.pdf>, Presentation at an intern symposium at MIT Lincoln Laboratory.